

# Data Link Layer

## Error Control: Introduction and Detection

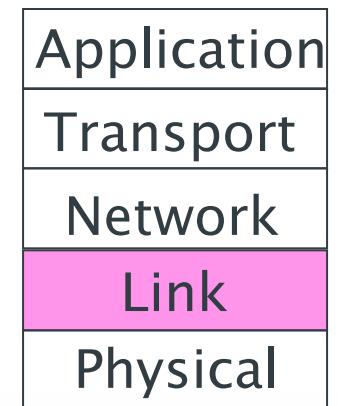
Geoff Merrett

ELEC3227/ELEC6255: Networks

*See Tanenbaum Chapter 3 (Data Link Layer)*

# Outline

- Error Control
  - Introduction
  - Error Detection
    - Parity (Single, Multiple, Interleaved)
    - Checksums
    - Cyclic Redundancy Checks (CRCs)
  - Error Correction
    - Hamming Codes



# Introduction

- Some transmission media are reasonably error free
  - e.g. a fibre optic link
- Some are not
  - e.g. a wireless link
- Error models
  - Single-bit errors
    - e.g. an extreme value of thermal noise
  - Bursts of errors
    - e.g. deep fading in a wireless channel;
    - e.g. transient electrical interference on a wired link etc

# Introduction

- We can use two strategies:
  - Error detection (using error detection codes), and then re-request frame
  - Error correction (using error correcting codes), aka Forward Error Correction (FEC)
- We should use something appropriate
  - If few errors, FEC will introduce unnecessary overhead
  - If many errors, retransmissions will introduce unnecessary overhead
- Might also be handled in other layers
  - FEC at the Physical and some higher layers (e.g. in real-time content streaming)
  - Error detection in Data Link, Network and Transport layers

# Redundancy

- We embed an ability to detect and/or correct errors by adding redundancy  
 $m$  (data bits) +  $r$  (redundant check bits) =  $n$  (total number of transmitted bits)
- This is then known as an  $(n, m)$  code with an  $n$ -bit codeword
- Only a fraction of the  $2^n$  possible codewords are used:  $2^m/2^n = 2^{-r}$ 
  - It is this sparseness or redundancy that allows errors to be detected/corrected.
- The code rate is the fraction of the codeword that carries non-redundant information, i.e. =  $m/n$ 
  - Noisy channel, a code rate of 0.5 might be suitable;
  - Reliable channel, a code rate close to 1.0 might be suitable.

# Block, systematic and linear codes

- Block code
  - Any set of input bits (the  $m$ -bit message) could be looked up in a table to find the corresponding  $n$ -bit codeword
    - i.e. the  $r$  check bits are solely a function of the  $m$ -bit message
- Systematic code
  - The set of input bits (the  $m$ -bit message) is transmitted alongside the  $r$  check bits
    - i.e. the  $m$ -bits in the message appear in the  $n$ -bits in the codeword
- Linear code
  - The check bits are a linear function of the message bits

# Hamming distance

- Hamming distance  $d$  = the number of bits different between two codewords
  - XOR and count the 1's
- If you know the algorithm used to create codewords, you can construct the complete set of all possible codewords.
  - $\min(d)$  is the Hamming distance of the complete code
- If two codewords are a Hamming distance  $d$  apart, it will require  $d$  single-bit errors to convert one into another
- Block codes can detect  $d - 1$  errors and correct  $(d - 1)/2$  errors

# Error Detection

# Single Parity Bit

- Even parity
  - add a redundant parity bit to make an even number of 1's
  - Equivalent of an XOR operation
  - E.g.  $1110000 \rightarrow 1110000\textcolor{magenta}{1}$
- Odd parity
  - add a redundant parity bit to make an odd number of 1's
- To check, see whether an even (or odd) number of bits are received

# Single Parity Bit

- If a single-bit error occurs:
  - e.g. 11100101; detected, sum is wrong
  - To be expected, as Hamming distance  $d = 1$ .
- If a double-bit error occurs:
  - e.g. 2 errors, 11101101; *not detected*, sum is correct!
- If any odd number of bit-errors occur:
  - e.g. 3 errors, 11011001; detected sum is wrong
- Can detect any odd number of errors (including in the parity bit itself)
  - i.e. random errors are detected with probability 0.5!

# Multiple Parity Bits

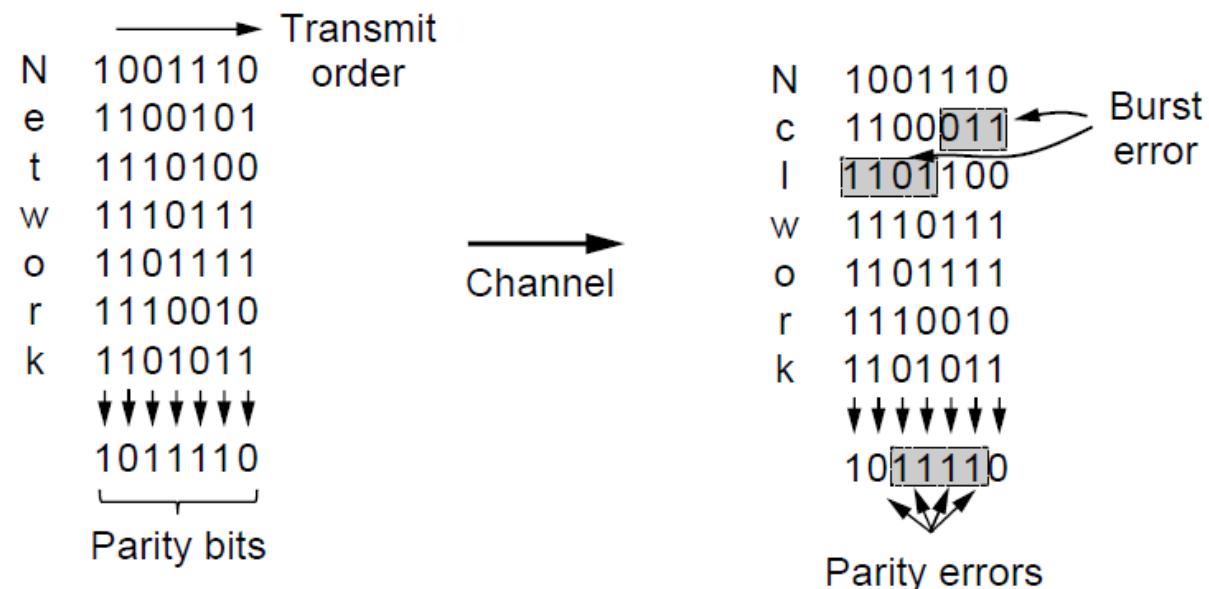
- Consider the message as a matrix  $w$  bits wide and  $h$  bits high
  - Send a separate parity bit for each  $w$ -bit row
  - $r$  (where  $r = h$ ) bit errors can be detected (provided max of 1 error occurs per row)
  - We add  $r$  parity bits e.g. for the message: 101100011000 (with  $w = 3$ ,  $h = 4$ , and even parity):

1	0	1	0
1	0	0	1
0	1	1	0
0	0	0	0

- Still not much good for burst errors (as maximum of 1 error per row!)

# Interleaved Parity

- Interleaving is a general technique
  - Converts a code that detects (or corrects) isolated errors...
  - ...into one that detects (or corrects) burst errors
- Interleaving  $r$  (where  $r = w$ ) parity bits detects burst errors up to length  $r$ 
  - Calculate parity bits over columns, but send data along rows
  - Parity bits are sent at the end
  - Each parity sum is made over non-adjacent bits
  - An even burst of up to  $r$  errors will not cause it to fail



# Checksums

- Often just used to refer to a set of check bits (e.g. parity bits)
- However, stronger checksums are based on summing the message bits
  - and are usually placed at the end of the message
- Typically treat the message as being formed from many  $N$ -bit words, and adds  $N$  (i.e.  $r = N$ ) check bits to the end which are the modulo  $2^N$  sum of all words
  - e.g. Internet 16-bit 1's complement checksum
- More effective than parity
  - e.g. if the LSB in two words have single bit flips, parity would not detect.
  - Detects bursts of up to  $r = N$  errors, and random errors with probability  $1-2^{-N}$
  - Vulnerable to systematic errors, e.g., added zeros, swapping parts of the message etc
- Other methods of creating a checksum can provide stronger protection
  - e.g. Fletcher's checksum: improves protection against changes in the position of data

# Cyclic Redundancy Checks (CRCs)

- Parity and checksums are rarely used in the Data Link Layer
  - A stronger error detection code is in widespread use: the CRC or polynomial code
  - e.g. can detect all double bit errors and not vulnerable to systematic errors
- Treats bit strings as the coefficients of a polynomial
  - A  $k$ -bit string is regarded as the coefficient of a polynomial with  $k$  terms
  - i.e.  $x^{k-1} + x^{k-2} + \dots + x^1 + x^0$  (a polynomial with a ‘degree’ =  $k-1$ )
- The protocol has an agreed generator polynomial,  $G(x)$ 
  - The degree of  $M(x)$  (the polynomial corresponding to the  $m$ -bit frame) must be greater than that of  $G(x)$

# Calculating CRCs

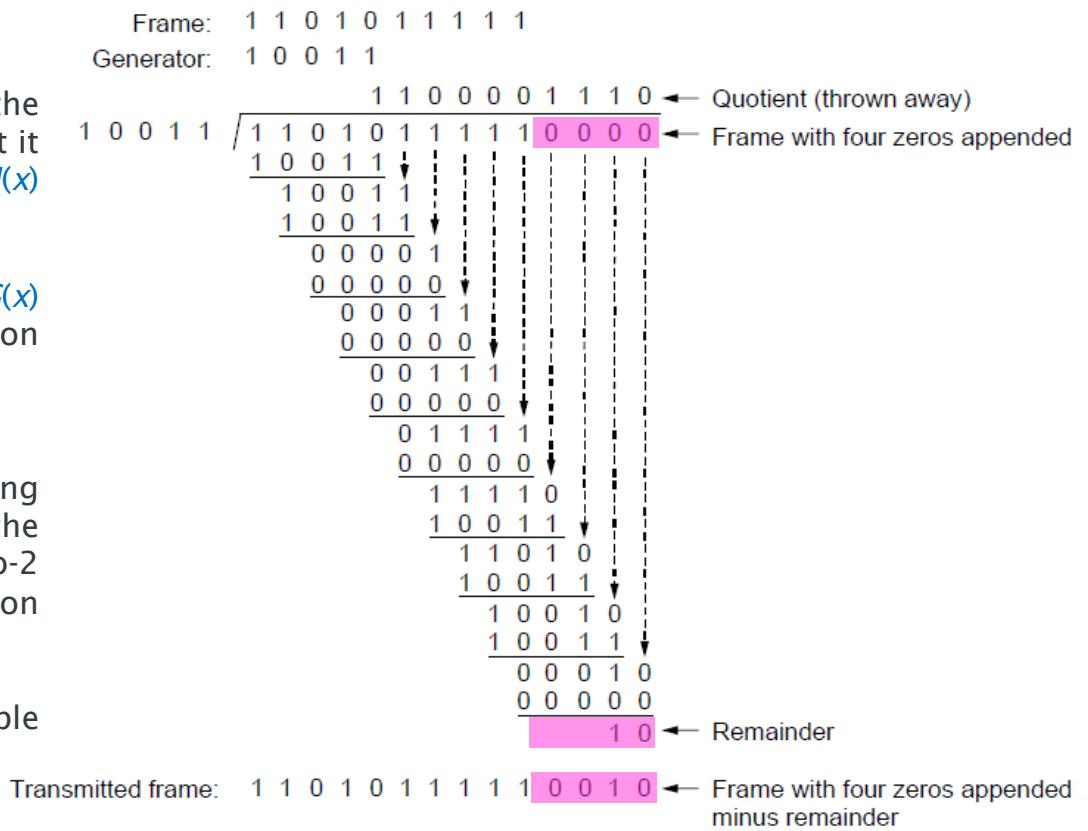
- Adds bits so that transmitted frame viewed as a polynomial is evenly divisible by a generator polynomial

Start by adding  $r (= k-1)$  0's to RHS of the  $M(x)$  bit string (the frame) – so that it becomes  $x^r \cdot M(x)$

Divide the bit string corresponding to  $G(x)$  into  $x^r \cdot M(x)$  using modulo 2 division

Create the  $n$ -bit bit string corresponding to  $T(x)$  for transmission by subtracting the remainder from  $x^r \cdot M(x)$  using modulo-2 subtraction

This makes  $T(x)$  evenly divisible



# Detecting Errors using CRCs

- $T(x)$  (the transmitted  $n$ -bit message) should clearly always divide exactly by  $G(x)$  as the remainder has been subtracted from it, e.g.
  - $123 / 10 = 12$  remainder 3
  - $(123 - 3) / 10 = 12$  remainder 0
- On decoding, all the receiver has to do is calculate  $T(x) / G(x)$ .
  - If the result has a remainder = 0, no errors were detected.
  - If errors have occurred during transmission, i.e.  $T(x) + E(x)$ , this will give a non-zero remainder when  $(T(x) + E(x)) / G(x)$  is calculated.
  - The remainder will only be zero if  $E(x)$  is a factor of  $G(x)$ 
    - *See book for more information on performance*
- CRCs are easy to calculate/check in hardware using shift and XOR operations

# IEEE 802 and CRC-32

- IEEE 802.3 (Ethernet) uses CRC-32
  - $x^{32} + x^{26} + x^{23} + x^{22} + x^{16} + x^{12} + x^{11} + x^{10} + x^8 + x^7 + x^5 + x^4 + x^2 + x^1 + 1$
  - Detects:
    - all burst errors of length 32 and less
    - all bursts affecting an odd number of bits
- IEEE 802.15.4 and LoRaWAN use the CCITT (ITU-T) CRC-16
  - $x^{16} + x^{12} + x^5 + 1$

# CRCs and Parity

- Even parity is the same as CRC-1
  - $G(x) = 1x^1 + 1x^0 = (x + 1)$
  - e.g. apply to the message:  $0110101_2$

# Error Correction

# Hamming Codes

- Consider a code that will allow all single-bit errors to be corrected
- Only a fraction of the  $2^n$  bit patterns (possible codewords) are used:  $2^m/2^n = 2^{-r}$ 
  - The others can be considered ‘illegal’ bit patterns
- Each of the  $2^m$  messages must have  $n$  illegal bit patterns one bit away from it
  - Identified by systematically inverting each of the  $n$ -bits in the valid codeword
- Therefore, each of the  $2^m$  messages requires  $n + 1$  bit patterns dedicated to it
  - Since there are a total of  $2^n$  bit patterns,  $(n + 1) \cdot 2^m \leq 2^n$
  - Because  $n = m + r$ , we can say that  $(m + r + 1) \leq 2^r$
  - This gives a theoretical lower-limit on  $r$ , the number of check bits needed to correct single errors
  - This theoretical limit can be achieved using Hamming codes

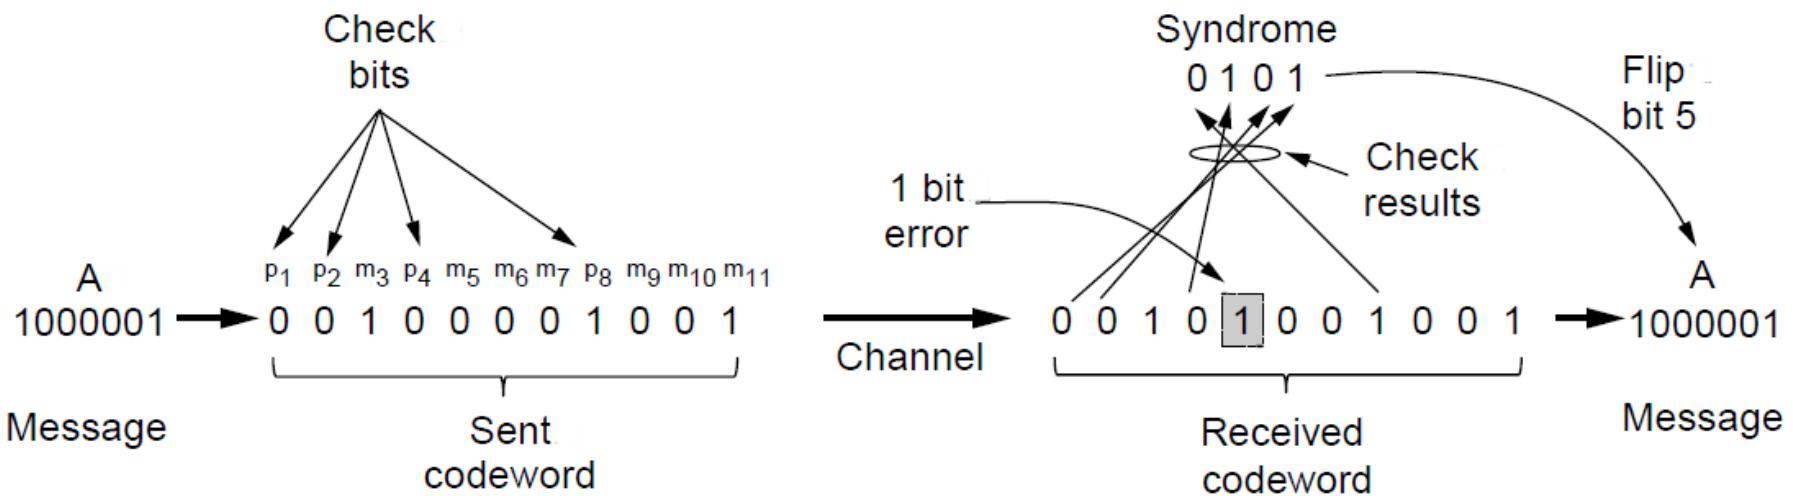
# Hamming Codes

- The  $n$  data bits are consecutively numbered from bit 1 at the left
- Bits in positions that are powers of 2 (1, 2, 4, 8, etc) are the  $r$  check bits
- The remaining bits are sequentially filled up with the  $m$  message bits
- The check bits are just even (or odd) parity bits
- Each of the  $m$  message bits may contribute to multiple check bits
  - Rewrite the location of a message bit as its powers of two, and those are the positions it contributes to
- Hamming codes provide a code with a Hamming distance  $d = 3$ 
  - An  $m=7$ -bit message, resulting in an  $n=11$ -bit codeword, is an  $(11, 7)$  Hamming code

# Hamming Codes

- Detecting errors
  - To ‘deconstruct’ the Hamming code, the check bits are recalculated at the destination (including the check-bit) – these are the *check results*
  - If even parity was used, the parity sum should be 0
  - If not, an error has been detected
- Correcting errors
  - If an error has been detected, the set of *check results* become the *error syndrome*
  - This indicates which bit was erroneous, and hence can simply be flipped

# Hamming Codes - example



(11, 7) Hamming code adds 4 check bits and can correct 1 error

# Hamming Codes vs Parity

- Consider a channel with a BER of  $10^{-6}$  and a block size of 1000 bits
- To correct a single-bit error (Hamming code):
  - We know that  $(m + r + 1) \leq 2^r$
  - Therefore,  $1001 \leq 2^r - r$ , and hence we'd need to add 10 check bits to each block
  - Therefore, a Mb of data would require an overhead of 10,000 check bits
- To detect a single-bit error (Parity)
  - We need only 1 check-bit per block
  - Therefore, a Mb of data would require an overhead of 1000 check bits
- Once every 1000 blocks, an error will occur
  - Error correction corrects the error: i.e. total overhead = 10,000 bits
  - Error detection retransmits the block: i.e total overhead =  $1000 + 1001 = 2001$  bits
- Therefore, in this case, error detection has a fifth of the overhead!

# ELEC3222 18/19 Exam Question

A 16-bit packet,  $1101\ 1111\ 1010\ 0110_2$ , is passed from the network layer to the data link layer for transmission.

1. The data link layer first applies interleaved even parity, with a width of 4 bits. **Calculate** the bit string produced.
2. **State** how many errors this can detect and correct. **Compare** this in terms of error detection/correction and overhead when compared to non-interleaved even parity (i.e. a parity bit after every 4 bits)?
3. The data link layer then frames the data using nibble (4-bit) stuffing. The flag and escape nibbles are  $0110_2$  and  $1111_2$  respectively. Using your answer from (1), **calculate** the bit string produced. If you were unable to answer (1), assume the bit string was unchanged.

# ELEC3227 19/20 Exam Question

A CRC is calculated for the bit string  $10011101_2$ , appended to the end, and transmitted. The generator  $G(x) = x^3 + 1$ .

- 1. Calculate** the actual transmitted bit string.
  
2. The most-significant-bit (MSB) is inverted during transmission. **Show** whether the receiver detects this and, if so, **explain** the steps it might take to rectify it.
  
3. **State** an example of a bit error in the transmitted bit string that will not be detected, and **explain** why.



Questions?